

## CSc 656 Project 1 Part b

(This document available on iLearn.)

due Monday 3/27/2017 3:00 pm (no grace period)

(4% of your grade)

This is an individual assignment. Work on your own!

Problem 4:

You are given two trace files (li1.trace and li2.trace) on unixlab.sfsu.edu in

~whsu/csc656/Traces/S17/P1

(These are in the unixlab file system. You can log on to unixlab.sfsu.edu, and access them through Unix file system commands. The directory is also linked to a web-accessible area: <http://unixlab.sfsu.edu/~whsu/csc656/Traces/S17/P1> )

Write two branch prediction simulators to work with these trace files. You may work in C/C++ or Java; other languages will require prior approval.

(There's also a third, very short trace file, test.trace, for preliminary debugging only.)

Each line of a trace file contains four hexadecimal integers, representing information from branches extracted from an instruction trace. The four (hexadecimal) integers on each line are:

- the address (i.e., PC) of the branch
- the *type* of the branch (1 means direct branch, meaning either a conditional or unconditional branch similar to a MIPS I-format branch; 2 means indirect branch, meaning a branch where the target address comes from a register, similar to the MIPS jr instruction)
- the target address
- whether the branch is taken (1) or not taken (0)

(These are SPARC traces, so a very small number of indirect branches are actually conditional.) For this project, your simulator will only work with conditional branches (type 1). **Discard all branches that are not of type 1.**

You will develop two branch prediction simulators, System 1 and System 2. **Your code must compile and run on unixlab.sfsu.edu.** Sample executables for these two systems, named sys1 and sys2, can be found on unixlab.sfsu.edu in ~whsu/csc656/CODE/S17/P1. Your programs will simulate the behavior of these branch predictors:

**System 1 (static prediction):**

Forward branches are predicted not taken, backward branches predicted taken. This is not a hardware branch predictor; you can get the misprediction statistics from scanning the list of conditional branches in each trace file. So if a forward branch is not taken, it is predicted correctly, otherwise it is mispredicted. Similarly, if a backward branch is taken, it is predicted correctly, otherwise it is mispredicted.

System 1 will take Unix command line arguments. The first argument is the name of the trace file. The second argument (optional) is “-v”, which turns on verbose mode. See **Submissions** section for format.

System 1 must work correctly for System 2 code to be graded. If your System 1 code does not work, don't even bother working on System 2.

### **System 2 (basic 2-bit predictor with branch target buffer):**

This is the N-entry 2-bit predictor/M-entry branch target buffer from Chapter 4a slides, in the section Branch Handling 2: dynamic branch prediction. Your code should work with any N and M, both powers of 2. Follow the definitions from the slides on how to index the array of 2-bit branch predictors, and how to update the states of each 2-bit predictor. Assume all predictors start in state 01. Remember that predictors are not tagged; different conditional branches can update the same predictor without reinitializing to state 01. The branch target buffer entries are tagged, of course.

System 2 operation in more detail:

```
For each branch  $B_i$ 
    Get prediction for  $B_i$  (predict taken or predict not taken)
    If  $B_i$  is predicted taken
        Check BTB entry for  $B_i$ 
        If valid bit of BTB entry == 1 && tag of BTB entry == tag of  $B_i$ 
            We have a BTB hit
        Else
            We have a BTB miss
    Else if  $B_i$  is predicted not taken
        Do nothing

    Check whether  $B_i$  is actually taken/not taken
    If prediction == actual behavior, prediction is correct
    Else prediction is wrong

    update prediction for  $B_i$ 
    If  $B_i$  actually taken
        Tag of BTB entry = tag of  $B_i$ 
        Valid bit of BTB entry = 1
```

System 2 will take Unix command line arguments. The first argument is the name of the trace file. The second argument is N (number of entries in predictor buffer, always a power of two), the third argument is M (number of entries in branch target buffer, always a power of two). The fourth argument (optional) is “-v”, which turns on verbose mode. See **Submissions** section for format.

### Simulator output

Sys1 should print

- the total number of conditional branches
- the number of forward branches
- the number of backward branches
- the number of forward taken branches
- the number of backward taken branches
- the number of mispredicted branches
- the misprediction rate for all branches ( $\# \text{ mispredictions} / \# \text{ branches}$ )

Sys2 should print all of the above, plus:

- the number of BTB misses
- the BTB miss rate ( $\# \text{ BTB misses} / \# \text{ BTB accesses}$ )

Your simulator should implement a verbose mode for debugging. Verbose mode is turned off by default, and turned on by the -v flag (see Submissions section for format). When verbose mode is on, in addition to the output generated in non-verbose mode, your simulator prints out, for each branch in the trace file, a list of integers on a single line. These are:

- Order of type 1 branch in trace file (the first type 1 branch in the file is numbered 0)
- Index of prediction buffer accessed (in hexadecimal)
- Current state of prediction buffer
- New state of prediction buffer after update
- Index of BTB accessed (in hexadecimal)
- Tag of BTB entry accessed (in hexadecimal)
- Number of BTB accesses so far
- Number of BTB misses so far

For example, the test trace *test.trace* contains:

```
18244 1 18338 1
1838c 1 18338 1
1838c 1 18338 1
1838c 1 18338 0
1a204 2 18210 1
18244 1 18338 1
1838c 1 18338 0
```

```
1a204 2 17fe8 1
18004 2 18380 1
1838c 1 18338 1
1838c 1 18338 0
1a204 2 17fe8 1
18004 1 1808c 1
```

The output of sys1 should be:

```
Number of branches = 9
Number of forward branches = 3
Number of forward taken branches = 3
Number of backward branches = 6
Number of backward taken branches = 3
Number of mispredictions = 6 0.666667
```

The verbose output of sys2 with 16 prediction buffers and 4 target buffers should be:

```
unixlab: ./sys2 test.trace 16 4 -v
0 1 1 2 1 1824 0 0
1 3 1 2 3 1838 0 0
2 3 2 3 3 1838 1 0
3 3 3 2 3 1838 2 0
4 1 2 3 1 1824 3 0
5 3 2 1 3 1838 4 0
6 3 1 2 3 1838 4 0
7 3 2 1 3 1838 5 0
8 1 3 3 1 1800 6 1
Number of branches = 9
Number of forward branches = 3
Number of forward taken branches = 3
Number of backward branches = 6
Number of backward taken branches = 3
Number of mispredictions = 6 0.666667
Number of BTB misses = 1 0.166667
unixlab:
```

## Measurements and results

You should make measurements for the two trace files, then fill out this table for each of the traces:

	# mispredictions	misprediction # BTB rate	BTB misses	BTB miss rate
System 1			N/A	N/A
System 2 (N=256,M=64)				
System 2 (N=1024,M=256)				

# mispredictions is the number of times the prediction from the prediction buffer does not match the behavior of the branch. Misprediction rate is # mispredictions / # branches.

A BTB miss occurs when a branch is predicted taken, but its target address is not in the target buffer. Branches that are not taken do not access the BTB. Hence, BTB miss rate is # BTB misses / # branches that are predicted taken.

Submit the result table with your code.

### Submission (source code + results table):

Submit a tar/zip file using the iLearn submission link. Your tar/zip file should expand into a single directory tagged with your name. The directory should contain your source files and the results table; each source file should have a header that with accurate instructions on compiling and running your code on the Unix command line. If your instructions don't work perfectly, you may get a zero on the project.

Your instructions should allow us to generate two executables, sys1 and sys2 (for System 1 and 2 respectively). We will then run each one with this command line (assuming we're on unixlab.sfsu.edu) for C/C++:

```
./sys1 ~whsu/csc656/Traces/S17/P1/li1.trace [-v]
./sys2 ~whsu/csc656/Traces/S17/P1/li2.trace 1024 256 [-v]
etc etc
```

Or for Java:

```
java sys1 ~whsu/csc656/Traces/S17/P1/li1.trace [-v]
java sys2 ~whsu/csc656/Traces/S17/P1/li2.trace 1024 256 [-v]
etc etc
```